

Details

Launch: bxcall\_cal\_flux\_residual\_kernel\_4 ( 9600, 1 1)Time: 2.30 msecondCycles: 1,945,424Registers: 168GPU: Quadro M5000SM Frequency: 845.40 MHzCC: 5.2

Current

Save as PDF

GPU Speed of Light

SOL SM [%]65.51Duration [msecond]2.30

SOL Memory [%]32.42Elapsed Cycles [cycle]19,45,424.75

SOL TEX [%]14.37SM Frequency [MHz]845.40

SOL L2 [%]32.42Memory Frequency [GHz]6.08

SOL FB [%]6.08

GPU Utilization

SM [%]65.51

Memory [%]32.42

Speed Of Light [%]

Recommendations

Bottleneck

[Warning] Compute is more heavily utilized than Memory. Look at "Compute Workload Analysis" report section to see what the compute pipelines are spending their time doing.  
Also, consider whether any computation is redundant and could be reduced.

Compute Workload Analysis

Executed Ipc Elapsed [inst/cycle]0.23SM Busy [%]65.51

Executed Ipc Active [inst/cycle]0.29Issue Slots Busy [%]4.87

Issued Ipc Active [inst/cycle]0.29

Memory Workload Analysis

Memory Throughput [Gbyte/second]11.05Mem Busy [%]32.42

L1 Hit Rate [%]31.76Max Bandwidth [%]31.69

L2 Hit Rate [%]90.81Mem Pipes Busy [%]1.68

Memory Chart

Kernel

Global325.45 k Inst9.60 k Req1.19 m Req

Local79.76 k Inst464.04 k Req33.60 k Req

Texture0.00 Inst0.00 Req

Surface0.00 Inst0.00 Req

Shared330.54 k Inst299.33 k Req338.37 k Req

Unified Cache

L2 Cache

System Memory

Device Memory

Shared Memory

	Instructions	Requests	% Peak	Bank Conflicts
Shared Load	1,58,604	2,99,328	1.24	28
Shared Store	1,71,940	3,38,374	1.40	0
Shared Atomic	0	-	-	-
Total	3,30,544	6,37,702	2.64	28

	Instructions	SM->TEX Requests	% Peak	Hit Rate	TEX->L2 Requests	% Peak	L2->TEX Returns	% Peak	TEX->SM Returns	% Peak
Global Load Cached	3,24,246	79,672	0.33	38.75	-	-	13,354	0.02	-	-
Global Load Uncached	-	11,06,304	4.57	-	-	-	84,64,629	13.60	-	-
Local Load Cached	75,564	4,800	0.02	-	-	-	4,800	0.01	-	-
Local Load Uncached	-	4,59,240	1.90	-	-	-	4,64,796	0.75	28,10,110	4.51
Surface Load	0	0	0	-	-	-	0	0	-	-
Texture Load	0	0	0	-	-	-	0	0	-	-
Global Store	1,200	9,600	0.04	-	38,400	0.06	-	-	-	-
Local Store	4,200	33,600	0.14	-	33,600	0.05	-	-	-	-
Surface Store	0	0	0	-	0	0	-	-	-	-
Global Reduction	0	0	0	-	0	0	-	-	-	-
Surface Reduction	0	0	0	-	0	0	-	-	-	-
Global Atomic	0	0	0	-	0	0	-	-	-	-
Global Atomic Cas	0	0	0	-	0	0	0	0	0	0
Surface Atomic	0	0	0	-	0	0	0	0	0	0
Surface Atomic Cas	0	0	0	-	0	0	0	0	0	0
Loads	3,99,810	16,50,016	6.82	0.09	-	-	89,47,579	14.37	28,10,110	4.51
Stores	5,400	43,200	0.18	-	72,000	0.12	-	-	-	-
Total	4,05,210	16,93,216	7.00	0.09	72,000	0.12	89,47,579	14.37	28,10,110	4.51

	TEX->L2 Requests	% Peak	L2->TEX Returns	% Peak	Total Bytes	Total Throughput
Global Load Cached	-	-	13,354	0.05	4,27,328	18,56,90,187.90
Global Load Uncached	-	-	84,64,629	29.98	27,08,68,128	1,17,70,81,57,192.12
Local Load Cached	-	-	4,800	0.02	1,53,600	6,67,48,247.86
Local Load Uncached	-	-	4,64,796	1.65	1,48,73,472	6,46,33,99,710.76
Surface Load	-	-	0	0	0	0
Texture Load	-	-	0	0	0	0
Global Store	38,400	0.14	-	-	12,28,800	53,39,85,982.87
Local Store	33,600	0	-	-	10,75,200	46,72,37,735.01
Surface Store	0	0	-	-	0	0
Global Reduction	0	0	-	-	0	0
Surface Reduction	0	0	-	-	0	0
Global Atomic	0	0	-	-	0	0
Global Atomic Cas	0	0	0	0	0	0
Surface Atomic	0	0	0	0	0	0
Surface Atomic Cas	0	0	0	0	0	0
Loads	-	-	89,47,579	31.69	28,63,72,528	1,24,42,40,04,338.64
Stores	72,000	0.14	-	-	23,04,000	1,00,12,23,717.88
Total	72,000	0.14	89,47,579	31.69	28,86,26,528	1,25,42,52,28,056.51

	L2<->FB Sectors	% Peak	Bytes	Throughput
Load	8,50,492	5.60	2,72,15,744	11,82,68,43,920.35
Store	72,465	0.48	23,18,880	1,00,76,89,954.39
Total	9,22,957	6.08	2,95,34,624	12,83,45,33,874.74

Scheduler Statistics

Active Warps Per Scheduler [warp/cycle]2.53Instructions Per Active Issue Slot [inst/issue]1.08

Eligible Warps Per Scheduler [warp/cycle]0.07No Eligible [%]92.21

Issued Warp Per Scheduler [issue/cycle]0.07One or More Eligible [%]6.97

Warps Per Scheduler

Theoretical Warps Per Scheduler

Active Warps Per Scheduler

Eligible Warps Per Scheduler

Issued Warp Per Scheduler

Recommendations

Issue Slot Utilization

[Warning] Every scheduler is capable of issuing two instructions per cycle, but for this kernel each scheduler only issues an instruction every 14.6 cycles. This might leave hardware resources underutilized and may lead to less optimal performance. Out of the maximum of 16 warps per scheduler, this kernel allocates an average of 2.53 active warps per scheduler, but only an average of 0.07 warps were eligible per cycle. Eligible warps are the subset of active warps that are ready to issue their next instruction. Every cycle with no eligible warp results in no instruction being issued and the issue slot remains unused. To increase the number of eligible warps either increase the number of active warps or reduce the time the active warps are stalled.

Warp State Statistics

Warp Cycles Per Issue Active [cycle/inst]34.16Avg. Active Threads Per Warp [thread/inst]28.14

Warp Cycles Per Issue Active [cycle/issue]36.83Avg. Not Predicated Off Threads Per Warp [thread/inst]27.62

Warp Cycles Per Executed Instruction [cycle/inst]34.18

Warp State (All Cycles)

Warp States

Stall Drain

Stall IMC Miss

Stall Long Scoreboard

Stall Barrier

Stall Member

Stall Short Scoreboard

Stall Tile Allocation

Stall Allocation Stall

Stall Wait

Stall No Instruction

Stall Math Pipe Throttle

Stall Tex Throttle

Stall MIO Throttle

Stall Dispatch Stall

Stall Misc

Stall Not Selected

Cycles per Instruction

Recommendations

CPI Stall 'Short Scoreboard'

[Warning] On average each warp of this kernel spends 24.7 cycles being stalled waiting for a scoreboard dependency on an MIO operation (not to TEX or L1). This represents about 67.1% of the total average of 36.8 cycles between issuing two instructions. The primary reason for a high number of stalls due to short scoreboards is typically memory operations to shared memory, but other contributors include frequent execution of special math instructions (e.g. MUFU) or dynamic branching (e.g. BRX, JMX). Consult the Memory Workload Analysis section to verify if there are shared memory operations and reduce bank conflicts, if reported.

CPI Stall 'Drain'

CPI Stall 'LG Throttle'

CPI Stall 'Long Scoreboard'

CPI Stall 'Math Pipe Throttle'

CPI Stall 'Member'

CPI Stall 'MIO Throttle'

CPI Stall 'Misc'

CPI Stall 'No Instructions'

CPI Stall 'Not Selected'

CPI Stall 'Sleeping'

CPI Stall 'TEX Throttle'

Instruction Statistics

Executed Instructions [inst]70,67,939Avg. Executed Instructions Per Scheduler [inst]1,10,436.55

Issued Instructions [inst]70,71,571Avg. Issued Instructions Per Scheduler [inst]1,10,493.30

Executed Instruction Mix

OpCodes

DFMA

DMUL

MOV

XMAD

MOV32i

IADD32i

IADD

LEA

DADD

LOP32i

STS

SYNC

SSY

LDS

ISETP

IADD3

DEPBAR

FSET

MUFU

BRA

LD

FFMA

DSETP

LOP

SHL

FSETP

SEL

IMNMX

ICADD

RET

CAL

BRK

PBK

F2F

SHR

I2F

LOP3

STL

S2R

STG

LDL

EXIT

Instructions Executed

Launch Statistics

Grid Size150Registers Per Thread [register/thread]168

Block Size64Static Shared Memory Per Block [Kbyte/block]12

Threads [thread]9,680Dynamic Shared Memory Per Block [byte/block]32

Waves Per SM1.56Shared Memory Configuration Size [Kbyte]10.81Block Limit SM [block]32

Block Durations

Warp Durations

Recommendations

Launch Configuration

Occupancy

Theoretical Occupancy [%]18.75Block Limit Registers [register]6

Achieved Occupancy [%]12Block Limit Local Mem [byte]8

Achieved Active Warps per SM [warp/cycle]15.64Block Limit Warps [warp]32

Achieved Active Warps Per SM [warp/cycle]10.81Block Limit SM [block]32

Impact of Varying Register Count Per Thread

Impact of Varying Block Size

Impact of Varying Shared Memory Usage Per Block